

Project 4

March 30, 2018

1 Project 04: Movielens Dataset Analysis

You don't need to limit yourself to the number of rows/cells provided. You can add additional rows in each section to add more lines of code.

Happy coding!

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sn
import re
%matplotlib inline
```

```
In [2]: df_movies=pd.read_csv('movies.dat',sep='::',names=['MovieID','Title','Genres'])
```

```
/home/ghanshyam/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: ParserWarning: Fa
""Entry point for launching an IPython kernel.
```

```
In [3]: def autolabel(rects):
    for rect in rects:
        h = rect.get_height()
        ax.text(rect.get_x()+rect.get_width()/2., 1.01*h, '%d'%int(h),
                ha='center', va='bottom')
```

```
In [4]: df_movies.head()
```

```
Out[4]:
```

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

```
In [5]: df_ratings=pd.read_csv('ratings.dat',sep='::',names=['UserID','MovieID','Rating','Timest
```

```
/home/ghanshyam/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: ParserWarning: Fa
""Entry point for launching an IPython kernel.
```

```
In [6]: df_ratings.head()
```

```
Out[6]:
```

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

```
In [7]: df_user=pd.read_csv('users.dat',sep=":",names='UserID::Gender::Age::Occupation::Zip-code')
```


```
/home/ghanshyam/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: ParserWarning: Failed to parse 'users.dat' as a CSV file.
    """Entry point for launching an IPython kernel.
```

```
In [8]: df_user.head()
```

```
Out[8]:
```

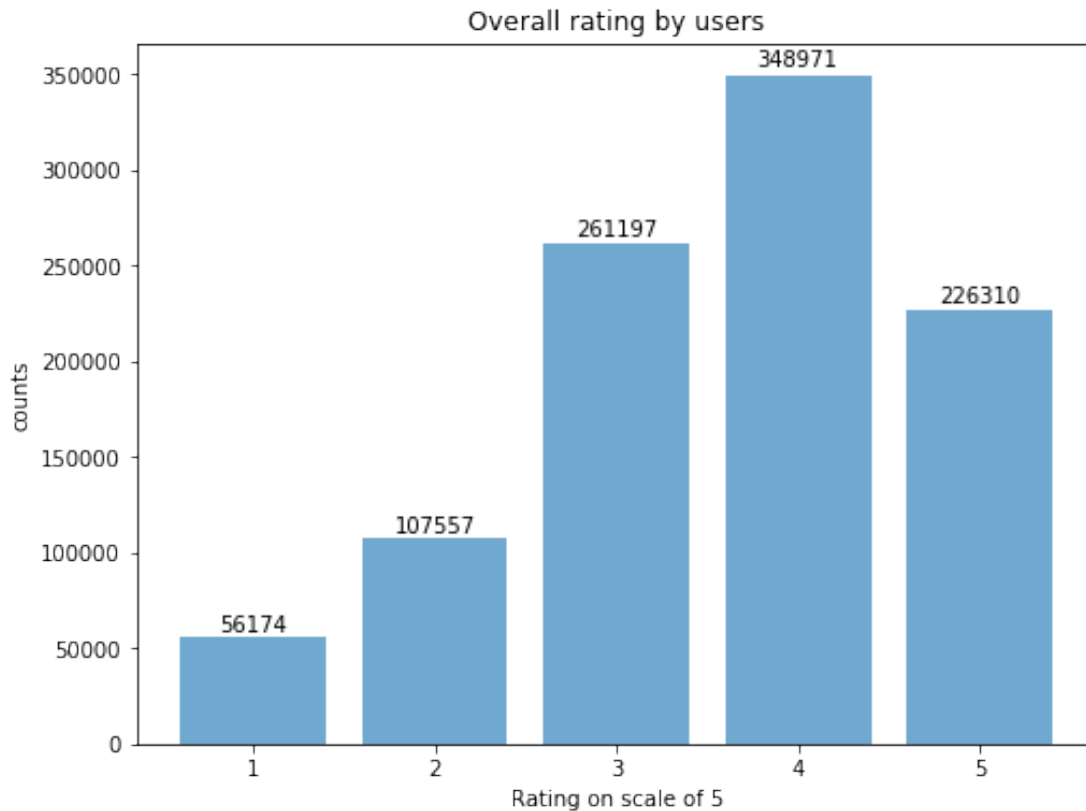
	UserID	Gender	Age	Occupation	Zip-code
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117
3	4	M	45	7	02460
4	5	M	25	20	55455

```
In [9]: age_dict={1:"1-18",
    18: "18-24",
    25: "25-34",
    35: "35-44",
    45: "45-49",
    50: "50-55",
    56: "56+"}
ageval=df_user.Age.apply(lambda x:age_dict[x] )
plt.figure(figsize=(8,6))
ageval.hist(bins=20,alpha=0.75)
plt.title('Age Distribution of users')
plt.xlabel('Range of age')
plt.ylabel('Number of users')
plt.show()
```



```
In [10]: fig=plt.figure(figsize=(8,6))
         ax=fig.add_subplot(111)
         s=df_ratings.Rating.value_counts()

         rect=ax.bar(s.index,s,alpha=0.64)
         plt.title('Overall rating by users')
         plt.xlabel('Rating on scale of 5')
         autolabel(rect)
         plt.ylabel('counts')
         plt.show()
```



2 Find and visualize the user rating of the movie “Toy Story”

```
In [11]: df_movies_toystory=df_movies[df_movies.Title.apply(lambda x: bool(re.search(r'^toy stor
```

```
In [12]: df_movies_toystory
```

```
Out[12]:
```

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
3045	3114	Toy Story 2 (1999)	Animation Children's Comedy

```
In [13]: fig = plt.figure(figsize=(12,8))
ax = fig.add_subplot(111)
bar_width=0.35
gap=0

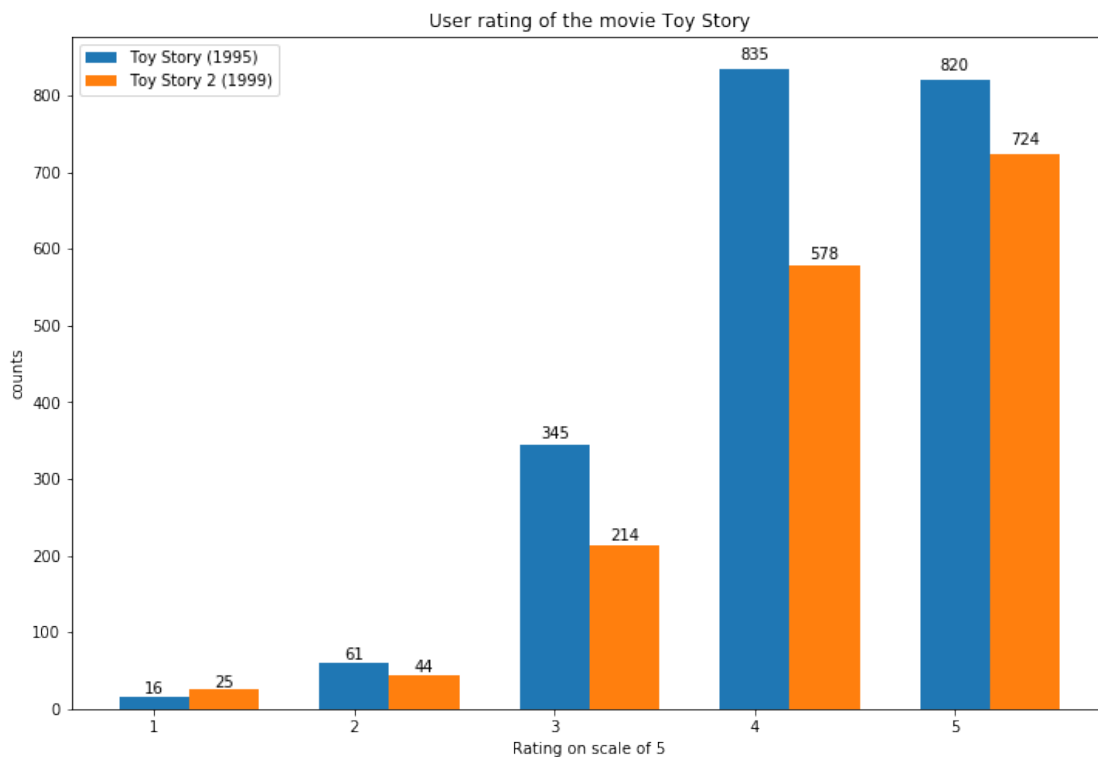
def autolabel(rects):
    for rect in rects:
        h = rect.get_height()
        ax.text(rect.get_x()+rect.get_width()/2., 1.01*h, '%d'%int(h),
                ha='center', va='bottom')
```

```

for movieId,title in zip(df_movies_toystory['MovieID'],df_movies_toystory['Title']):
    df_tmp_ratings=df_ratings[df_ratings['MovieID']==movieId]
    s=df_tmp_ratings.Rating.value_counts()
    #p=s.index.tolist()
    #print('s=',s,'\n dddindex=',p)
    rect=ax.bar(s.index+gap,s,bar_width,label=title)
    autolabel(rect)

    gap+=bar_width
plt.title('User rating of the movie Toy Story')
plt.xlabel('Rating on scale of 5')
plt.ylabel('counts')
plt.legend()
plt.show()

```



3 Find and visualize the viewership of the movie “Toy Story” by age group

```

In [14]: t1=df_movies_toystory.set_index('MovieID').join(df_ratings.set_index('MovieID'))
         t1['MovieID']=t1.index
         df_movie_age=t1.set_index('UserID').join(df_user.set_index('UserID'))
         df_movie_age.head()

```

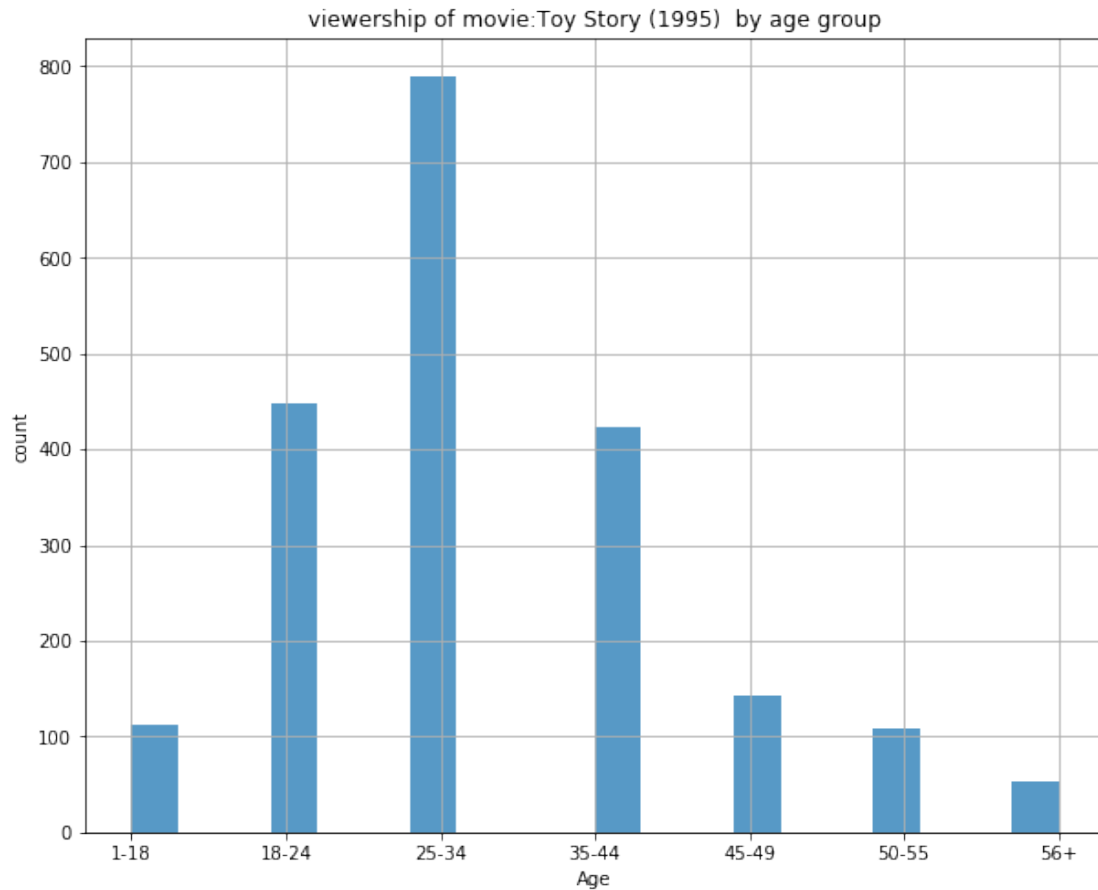
```
Out[14]:
```

	Title	Genres	Rating	Timestamp	\
UserID					
1	Toy Story (1995)	Animation Children's Comedy	5	978824268	
1	Toy Story 2 (1999)	Animation Children's Comedy	4	978302174	
3	Toy Story 2 (1999)	Animation Children's Comedy	3	978298103	
6	Toy Story (1995)	Animation Children's Comedy	4	978237008	
8	Toy Story (1995)	Animation Children's Comedy	4	978233496	


	MovieID	Gender	Age	Occupation	Zip-code
UserID					
1	1	F	1	10	48067
1	3114	F	1	10	48067
3	3114	M	25	15	55117
6	1	F	50	9	55117
8	1	M	25	12	11413

```
In [15]: fig=plt.figure(figsize=(10,8))
ageformovieid_1=df_movie_age.Age[df_movie_age['MovieID']==1]
ageformovieid_1=ageformovieid_1.apply(lambda x:age_dict[x])
ageformovieid_1.hist(bins=20,alpha=0.75)

plt.title('viewership of movie:Toy Story (1995) by age group')
plt.ylabel('count')
plt.xlabel('Age')
plt.show()
```



```
In [16]: plt.figure(figsize=(10,8))
ageformovieid_1=df_movie_age.Age[df_movie_age['MovieID']==3114]
ageformovieid_1=ageformovieid_1.apply(lambda x:age_dict[x])
ageformovieid_1.hist(bins=20,alpha=0.75)
plt.title('Viewership of movie:Toy Story 2 (1999) by age group')
plt.ylabel('count')
plt.xlabel('Age')
plt.show()
```



4 Find and visualize the top 25 movies by viewership rating

```
In [17]: df_ratings.head()
```

```
Out[17]:   UserID  MovieID  Rating  Timestamp
0        1      1193        5   978300760
1        1        661        3   978302109
2        1        914        3   978301968
3        1      3408        4   978300275
4        1      2355        5   978824291
```

```
In [18]: df_movie_rating=df_ratings[['MovieID','Rating']]
top25=df_movie_rating.groupby(['MovieID']).mean()['Rating'].nlargest(25)
top25=pd.DataFrame(top25)
top25.index
```

```
Out[18]: Int64Index([ 787,  989, 1830, 3172, 3233, 3280, 3382, 3607, 3656, 3881, 3245,
                    53, 2503, 2905, 2019,  318,  858,  745,   50,  527, 1148,  439,
                    557,  578, 1795],
                    dtype='int64', name='MovieID')
```



```
In [19]: df_rating_25=df_ratings[df_ratings['MovieID'].isin(top25.index)]
df_rating_25=df_rating_25[['MovieID','Rating']]
print('filtered',df_rating_25.shape,'from ',df_ratings.shape)
```

filtered (10816, 2) from (1000209, 4)

```
In [20]: tmp=df_rating_25.groupby(['MovieID','Rating'])['Rating'].count()
tmp.head()
```

```
Out[20]: MovieID  Rating
50           1         8
           2        31
           3       136
           4       464
           5      1144
Name: Rating, dtype: int64
```

```
In [21]: np_movieid_rating=np.array([list(t) for t in tmp.index])
df_heatmap=pd.DataFrame({})
df_heatmap['rating count']=tmp
df_heatmap['MovieID']=np_movieid_rating[:,0]
df_heatmap['rating']=np_movieid_rating[:,1]
df_heatmap.head()
```


```
Out[21]:
```

		rating count	MovieID	rating
MovieID	Rating			
50	1	8	50	1
	2	31	50	2
	3	136	50	3
	4	464	50	4
	5	1144	50	5

```
In [22]: hitmap_pivot=df_heatmap.pivot(index='MovieID',columns='rating',values='rating count')
```

```
In [23]: hitmap_pivot=hitmap_pivot.fillna(0)
plt.figure(figsize=(15,8))
sn.heatmap(hitmap_pivot,annot=True)
plt.title('Count of rating for top 25 movies')
plt.plot()
```

```
Out[23]: []
```



5 Find the rating for a particular user of user id = 2696

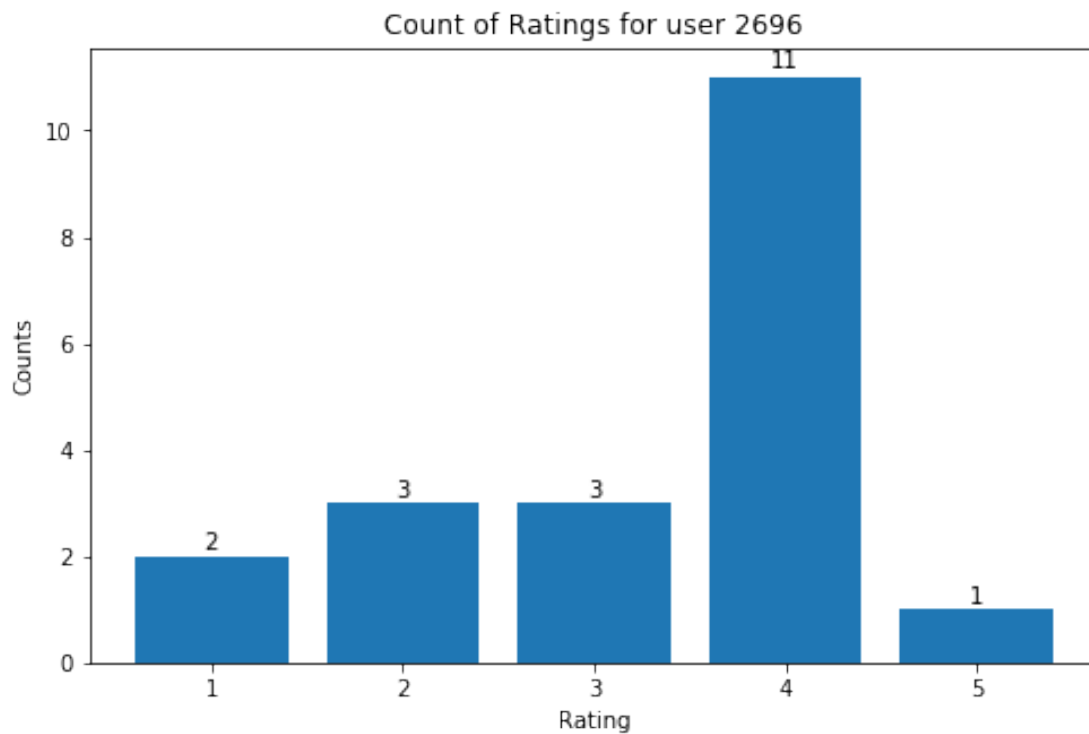
```
In [24]: df_ratings_2696=df_ratings[df_ratings['UserID']==2696]
df_ratings_2696
```

```
Out[24]:
```

	UserID	MovieID	Rating	Timestamp
440667	2696	1258	4	973308710
440668	2696	1270	2	973308676
440669	2696	1617	4	973308842
440670	2696	1625	4	973308842
440671	2696	1644	2	973308920
440672	2696	1645	4	973308904
440673	2696	1805	4	973308886
440674	2696	1892	4	973308904
440675	2696	800	5	973308842
440676	2696	2338	2	973308920
440677	2696	1711	4	973308904
440678	2696	3176	4	973308865
440679	2696	2389	4	973308710
440680	2696	1589	3	973308865
440681	2696	2713	1	973308710
440682	2696	3386	1	973308842
440683	2696	1783	4	973308865
440684	2696	350	3	973308886

440685	2696	1092	4	973308886
440686	2696	1097	3	973308690

```
In [25]: fig=plt.figure(figsize=(8,5))
ax=fig.add_subplot(1,1,1)
t=df_ratings_2696[['MovieID','Rating']]
p=t['Rating'].value_counts()
rect=ax.bar(p.index,p)
autolabel(rect)
plt.title('Count of Ratings for user 2696')
plt.xlabel('Rating')
plt.ylabel('Counts')
plt.show()
```



```
In [26]: df_ratings_2696.head()
```

```
Out[26]:
```

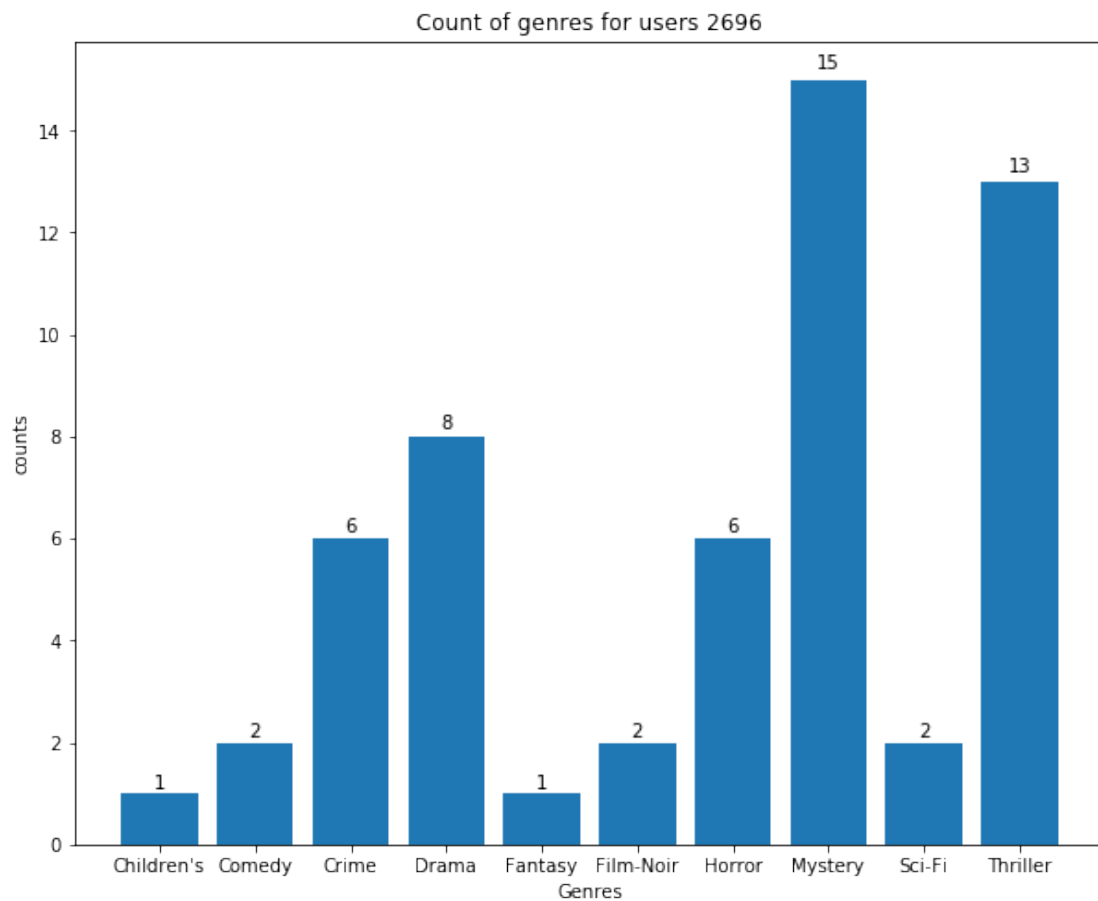
	UserID	MovieID	Rating	Timestamp
440667	2696	1258	4	973308710
440668	2696	1270	2	973308676
440669	2696	1617	4	973308842
440670	2696	1625	4	973308842
440671	2696	1644	2	973308920

```
In [27]: t=df_movies[df_movies['MovieID'].isin(df_ratings_2696['MovieID'])]
Geners=[]
```

```

for i in t['Genres']:
    for g in i.split('|'):
        Geners.append(g)
Geners=pd.Series(np.array(Geners))
p=Geners.value_counts()
fig=plt.figure(figsize=(10,8))
ax=fig.add_subplot(111)
rect=ax.bar(p.index,p)
autolabel(rect)
plt.title('Count of genres for users 2696')
plt.xlabel('Genres')
plt.ylabel('counts')
plt.show()

```



6 Machine Learning

```

In [98]: from sklearn.neighbors import KNeighborsClassifier
         from sklearn.linear_model import LogisticRegression

```

```

from sklearn.svm import SVC
import sklearn.cross_validation as cv
from sklearn import tree

```

```

In [99]: model_knn=KNeighborsClassifier(n_neighbors=3)
        model_log=LogisticRegression()
        model_svm=SVC(probability=True)
        model_decisiontree=tree.DecisionTreeClassifier(max_depth=5)

```

prepare data

```

In [100]: df_ratings_500=df_ratings.head(500)

```

```

In [101]: #movie id, age, occupation
        df_ratings_500['age']=df_ratings_500['UserID'].map(df_user.set_index('UserID')['Age'])
        df_ratings_500['occupation']=df_ratings_500['UserID'].map(df_user.set_index('UserID')['occupation'])
        df_ratings_500.head()
        #df_ratings_500.to_csv('df_rating_500.csv')

```

/home/ghanshyam/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

/home/ghanshyam/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>
This is separate from the ipykernel package so we can avoid doing imports until

```

Out[101]:
   UserID  MovieID  Rating  Timestamp  age  occupation
0        1     1193        5   978300760    1          10
1        1      661        3   978302109    1          10
2        1      914        3   978301968    1          10
3        1     3408        4   978300275    1          10
4        1     2355        5   978824291    1          10

```

```

In [102]: x_feature=df_ratings_500[['MovieID','age','occupation']]
        y_target=df_ratings_500[['Rating']]

```

```

In [103]: print(x_feature.shape,y_target.shape)

```

```

(500, 3) (500, 1)

```

```

In [104]: x_feature=np.array(x_feature)
        y_target=np.array(y_target).ravel()

```

```
In [105]: x_train,x_test,y_train,y_test=cv.train_test_split(x_feature,y_target)
```

```
In [106]: model_knn.fit(x_train,y_train)
          model_log.fit(x_train,y_train)
          model_svm.fit(x_train,y_train)
          model_decisioontree.fit(x_train,y_train)
```

```
Out[106]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=5,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

```
In [107]: score_knn=model_knn.score(x_test,y_test)
          score_log=model_log.score(x_test,y_test)
          score_svm=model_svm.score(x_test,y_test)
          score_decision_tree=model_decisioontree.score(x_test,y_test)
          print('score_knn: ',score_knn,'score_log: ',score_log,'score_svm:',score_svm,'decision
```

```
score_knn:  0.184 score_log:  0.4 score_svm: 0.352 decision tree: 0.336
```